

# Hamid Rezaee

+1 (607) 663-1415 | [hr328@cornell.edu](mailto:hr328@cornell.edu) | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | Ithaca, NY

## EDUCATION

---

### Cornell University

Ithaca, NY, USA

Bachelor's, Information Science

Expected May 2026

- Relevant Coursework: Object-Oriented Programming & DSA (Java), Programming with Python, Reinforcement Learning, Backend Development, Machine Learning for Business, Discrete Structures, Statistical Analysis, Advanced NLP, Advanced Information Retrieval, Machine Learning Independent Research, Technical Project Management

## EXPERIENCES

---

### Pequity

Ithaca, NY, USA

Machine Learning Engineering Intern

August 2024 – Present

- Built Scout, an interactive IR application using Node.js and React.js for retrieving and matching job titles to skills
- Accelerated an API-based data retrieval process by 3.6 times through multi-threaded implementation in Python
- Implemented agentic compensation pipeline, cutting down research delay from 12 minutes/employee to 20 seconds
- Executed semantic clustering using Hugging Face's ModernBERT for memory efficiency and clustering accuracy

### Dalus, Inc. (YC W25)

San Francisco, CA, USA

Software Engineering Intern (AI/ML Engineering)

June 2025 – August 2025

- Built custom LSA retrieval engine replacing Gemini 2.5 Flash, saving \$24k annually and reducing latency 10s → 50ms
- Architected schema-agnostic JSON Patch system for AI agent modifications, eliminating runtime errors in production
- Deployed GPU-powered document processing on NVIDIA T4, cutting latency 20s → 70ms and saving \$9K annually
- Engineered dynamic context management for AI copilot, solving overflow issues to enable stateful conversation

### Cornell Tech

New York City, NY, USA

Deep Learning Engineering Intern

May 2024 – May 2025

- Developed a loss function for small lesion segmentation using PyTorch and trained a model on NVIDIA 2xA6000 GPUs
- Beat the 81% benchmark by achieving an 83% mean accuracy score on validation split of the small lesion dataset
- Trained SliceSense for predicting MRI slice thickness using PyTorch with ResNet-BiLSTM-Attention architecture
- Delivered 96.5% accuracy on slice thickness prediction for MRI and abdominal imaging over 6 thickness classes

### Cornell University Sustainable Design

Ithaca, NY, USA

Machine Learning Engineering Lead

August 2023 – August 2024

- Led a group of five students to develop a model predicting HVAC components and leading to energy efficiency
- Trained XGBoost and LightGBM models and optimized it for inference with successful backend integration using Flask
- Execute rigorous feature engineering leading to 18% model dimensionality reduction and inference acceleration
- Managed cross-team development efforts with the backend team by organizing and leading weekly meetings

## PROJECTS

---

### ReSearch

Ithaca, NY, USA

Lead Developer and Project Manager

February 2025 – April 2025

- Led a team of four junior developers and programmed the full-stack application using JavaScript D3, Flask, and Python
- Developed SVD-based retrieval and elegant interactive visual design for frontend to enhance user interaction quality

### SliceSense

Ithaca, NY, USA

Deep Learning Researcher

March 2025 - April 2025

- Trained and optimized ResNet-BiLSTM using PyTorch and Torchvisions's ResNet-18 for training and inference efficiency
- Achieved a 96% post-TTA test-set accuracy on abdominal data and a 97% raw test-set accuracy on MRI data

### Alignment Arena

Ithaca, NY, USA

Machine Learning Researcher

January 2025 - April 2025

- Utilized Mixtral-8x7B and Hugging Face Transformers to run inference on NVIDIA A100 GPU and generate synthetic data
- Performed bias quantification over five Hugging Face masked language models using HF's Transformers and PyTorch

## SKILLS

---

**Programming Language:** Python, JavaScript, TypeScript, R, SQL, Java

**Platforms and Libraries:** PyTorch, TensorFlow, JAX, Scikit-learn, LightGBM, XGBoost, Transformers, Hugging Face, LangChain, AWS Bedrock, SageMaker, CloudFormation, Lambda, S3, EC2, ECS, Node.js, React.js, Next.js, D3.js, Tailwind, Git, Docker, Vercel, CUDA, RAG, LSTM, K-Means, Advanced NLP, Large Language Models (LLMs)